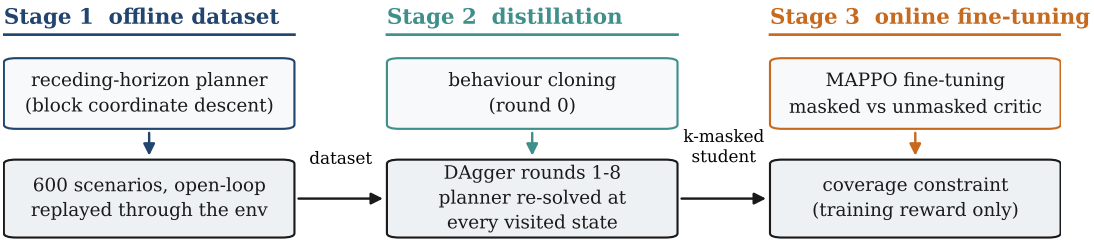
Stage 1 offline dataset
Stage 2 distillation
Stage 3 online fine-tuning
receding-horizon planner
(block coordinate descent)
behaviour cloning
(round 0)
MAPPO fine-tuning
masked vs unmasked critic
600 scenarios, open-loop
replayed through the env
dataset
DAgger rounds 1-8
planner re-solved at
every visited state
k-masked
student
coverage constraint
(training reward only)
sample budget (team steps)
distilled pipeline 341k (offline 53k + online 288k)
from scratch, to the same level 1.13M (3.3x)